

Resilient Next-Hop Groups in Linux

Ido Schimmel, Petr Machata

Nvidia

idosch@nvidia.org, petrm@nvidia.org

Abstract

Multipath next-hop groups are used to describe a situation in which it is possible to forward packets through one of several next hops, each of which will bring the packet in question closer to the destination. A typical algorithm to decide which next-hop to forward the packet through is hash-threshold, whereby a hash derived from packet headers is compared against hash space ranges assigned to individual next hops. The downside of this algorithm is that next hop addition, removal, and weight adjustment may end up redirecting flows that have previously been forwarded to unrelated next hops. This in turn may cause connection resets and performance hits. This paper presents resilient next-hop groups, an approach to hash space management that minimizes incidental hash space reassignment across next hop additions, removals, and weight adjustments.

Keywords

Linux, ECMP, next-hop groups, resilient hashing

Motivation

ECMP is typically used in one of two scenarios. To balance traffic between different paths that all lead to the same destination. Or to balance traffic between different servers with the same anycast IP address. This is illustrated by figure 1.

Common strategies when deciding which of the available next hops to forward the packet through include modulo-N and hash-threshold algorithms[1]. In both cases, hash over several packet fields is first obtained. In the case of modulo-N algorithm, the next hop to choose is decided simply by applying a modulo-N operation (with N the number of paths) on the packet hash. In the case of hash-threshold, each next hop is assigned a contiguous area of the hash space. A packet is then forwarded through the next hop that contains the packet's hash.

When number of next hops changes, both algorithms change the next hop that packets with certain hashes would be forwarded to. The amount of this disruption depends on the algorithm. Under modulo-N, the mapping from hashes to next hop is just completely different after the number of next hops changes. Hash-thresholds mitigates this disruption, but is not immune to it either. Figure 2 illustrates what happens as one next hop is removed.




Figure 2: As next hop C is removed, a part of hash space previously assigned to B is reassigned to A, and parts of D to E.

The disruption outlined above in practice means that certain traffic flows may get redirected through a different path or to a different server. The former may lead to packet reordering, which might have performance impacts. The latter is problematic since an established TCP connection forwarded to a server that is not familiar with the connection will result in the connection being reset.

Approach

The core idea behind resilient next-hop groups is to split the hash space into regular-sized, but fine-grained buckets, and assign these to next hops arbitrarily.

By permitting this fine-grained assignment, it is then possible to reassign only those parts of hash space that it is actually necessary to reassign, or whose reassignment causes the least disruption. If the granularity is sufficiently fine it is still possible to model different counts of next hops with varying weights without introducing errors. And the fact that the buckets are regular makes the solution manageable in terms of implementation effort, runtime performance and possible in-HW implementation.

Next-hop selection algorithm then again becomes modulo-N, except now N refers to the number of buckets. The next hop to forward through is the one that this bucket is assigned





Figure 1: ECMP use cases. Left: balancing between different paths leading to the same server. Right: balancing between servers.

to. Since the number of buckets is fixed, the choice of the modulo-N algorithm will not negatively impact the amount of traffic disruptions. See figure 3 for illustration.




Figure 3: Resilient next-hop groups introduce a layer of indirection between the hash and the next hops.

A case where this algorithm works very well is next-hop removal. As a next hop is removed, any flows hitting that next hop will certainly be disrupted anyway. Therefore the buckets assigned to the removed next hop can be freely distributed among the existing next hops according to their weights. There is no incidental disruption, as the traffic hitting the buckets assigned to the other next hops keeps getting resolved to the same next hops as before. This is illustrated in figure 4.




Figure 4: As next hop C is removed, its buckets are reassigned to other next hops. There is no incidental disruption.

Next hop addition is a bit more difficult. The issue is that

hash space for the new next hop has to come from somewhere, and there is no other way to get it but to take from the current next hops. (This is illustrated in figure 5.) The algorithm therefore has no choice but to disrupt some flows. This is unlike the removal case, where the choice of which flows to disrupt was imposed by external reality.




Figure 5: As next hop F is added, some buckets must be re-assigned from other next hops to satisfy F's hash space demands.

The resilient next-hop algorithm has three-tiered approach to minimizing this disruption.

The first tier is keeping track of bucket activity. The idea here is that each bucket knows when it was last used. A system administrator, as they are configuring the next-hop group, decides after how much time of not seeing traffic a given bucket is considered idle. Then when choosing which buckets to reassign, only idle buckets are eligible.

However bucket activity depends on traffic being forwarded, and a particularly unfortunate pattern could mean that there are no idle buckets that it makes sense to reassign. That is when the second tier comes into play. In that situation the algorithm will back off for a while, and will attempt to reassign buckets later.

Obviously the situation may be such that there keep being no idle buckets to reassign. The third tier is a possibility to force-balance the next-hop group. In that case the resilient group can be configured to reassign buckets without regard to their idleness.

Algorithm

In a nutshell, the algorithm works as follows. Each next hop deserves a certain number of buckets, according to its weight and the number of buckets in the hash table. In accordance with the source code, we will call this number a "wants count" of a next hop. In case of an event that might cause bucket allocation change, the wants counts for individual next hops are updated.

Next hops that have fewer buckets than their wants count, are called "underweight". Those that have more are "overweight". If there are no overweight (and therefore no underweight) next hops in the group, it is said to be "balanced".

Each bucket maintains a last-used timer. Every time a packet is forwarded through a bucket, this timer is updated to current jiffies value. One attribute of a resilient group is then the "idle timer", which is the amount of time that a bucket must not be hit by traffic in order for it to be considered "idle". Buckets that are not idle are busy.

After assigning wants counts to next hops, an "upkeep" algorithm runs. For buckets:

- that have no assigned next hop, or
- whose next hop has been removed, or
- that are idle and their next hop is overweight,

upkeep changes the next hop that the bucket references to one of the underweight next hops. If, after considering all buckets in this manner, there are still underweight next hops, another upkeep run is scheduled to a future time.

There may not be enough "idle" buckets to satisfy the updated wants counts of all next hops. Another attribute of a resilient group is the "unbalanced timer". This timer can be set to 0, in which case the table will stay out of balance until idle buckets do appear, possibly never. If set to a non-zero value, the value represents the period of time that the table is permitted to stay out of balance.

With this in mind, we update the above list of conditions with one more item. Thus buckets:

- whose next hop is overweight, and the amount of time that the table has been out of balance exceeds the unbalanced timer, if that is non-zero,
- ... are migrated as well.

Implementation

In Linux 5.3, support for standalone next-hop objects has been added to the Linux kernel. Resilient next-hop groups have been implemented as a new next-hop group type.

User Interface

The fact that resilient next-hop groups are implemented on top of the standalone next-hop objects informs how the `iproute2` CLI will look like. A new `type` keyword has been added to the `nexthop` group suite of commands. When `type` is `resilient`, a number of parameters to configure the resilient-specific attributes of the next-hop group can be specified. For example:

```
# ip nexthop add id 1 via 192.0.2.2 dev dummy1
# ip nexthop add id 2 via 198.51.100.2 dev dummy2
# ip nexthop add id 10 group 1/2 type resilient \
    buckets 8 idle_timer 120 unbalanced_timer 300
```

... first creates two next hops, and then groups them in a resilient group. `buckets` refers to number of buckets (and therefore hash space assignment granularity). The size of 8 indicated here is probably too low and is used for illustration purposes only. Typically one wants approximately hundreds of buckets, so that different counts of next hops and various odd weight ratios can be modeled accurately.

`idle_timer` and `unbalanced_timer` are discussed above.

Individual parameters (except for the number of bucket, which is fixed) can be changed. E.g. in the following, only, respectively, `idle_timer` and `unbalanced_timer` are adjusted, with the other staying intact:

```
Change attributes of the group
# ip nexthop replace id 10 group 1/2 \
    type resilient idle_timer 100
# ip nexthop replace id 10 group 1/2 \
    type resilient unbalanced_timer 900
```

Of particular interest is then that in neither of these cases, any bucket reassignments are done. The bucket table stays intact. Likewise in the following:

```
# ip nexthop replace id 10 \
    group 1,9/2,11 type resilient
```

Here we have changed the weights of the next hops. But as explained above, that does not by itself change any bucket reassignments. Instead the group become unbalanced, because next hop 1 now has more buckets than it should have, and next hop 2 has fewer. Then the upkeep algorithm gradually brings the group into balance (or not, depending on the configuration).

It is obviously possible to dump the group as well:

```
# ip nexthop show id 10
id 10 group 1/2 type resilient
buckets 8 idle_timer 60
unbalanced_timer 300 unbalanced_time 0
```

The `unbalanced_time` (as opposed to `unbalanced_timer`) shows the amount of time that the group has been out of balance. The 0 indicated here means that the group is actually balanced.

Finally, the netlink API has been extended with a new message type, `RTM_GETNEXTHOPBUCKET`, to dump individual buckets. That API is implemented in `iproute2`, too, and thus it is possible to introspect the state of the bucket table:

```
# ip nexthop bucket show id 10
id 10 index 0 idle_time 5.59 nhid 1
id 10 index 1 idle_time 5.59 nhid 1
id 10 index 2 idle_time 8.74 nhid 1
id 10 index 3 idle_time 8.74 nhid 1
id 10 index 4 idle_time 8.74 nhid 2
id 10 index 5 idle_time 8.74 nhid 2
id 10 index 6 idle_time 8.74 nhid 2
id 10 index 7 idle_time 8.74 nhid 2
```

(And now you know why the group only has 8 buckets.)

Offloading

Resilient next-hop groups are not novel, in fact the algorithm is commonly implemented in SDKs for networking switches and routers. The fact that it has not been implemented in the software datapath so far perhaps indicates that it is less critical in software-only deployments. It is therefore important to make it possible to implement offloading of resilient next-hop groups.

The resilient algorithm hinges on the fact that it understands which next hop buckets are idle. In software datapath, this understanding is achieved by maintaining bucket activity. However, when accelerating traffic by offloading to capable devices, the vast majority of traffic is never seen by the algorithm.

A two-pronged approach is implemented to close this loop-hole.

First, as is usual, the core next-hop algorithm communicates information about the events in resilient next-hop groups using the in-kernel notification mechanism. Bucket migration notifications in particular are of two flavors: forced and vetoable. Forced notifications are used if a next hop is removed and the algorithm simply needs to migrate a given bucket to another next hop. Vetoable notifications are used for the regular table upkeep, as the algorithm finds new idle buckets and proposes them for migration. The idea is that a driver will take note of these proposals, and in turn will propose them to the HW, which may bounce them if hardware datapath traffic has been using the bucket in question.

The second prong is a proactive reporting by the driver towards the next-hop code. A new in-kernel API `nexthop_res_grp_activity_update()` can be called to feed an activity bit-vector from the device to the next-hop group. Set bits in the vector are treated as if software-datapath traffic hit the corresponding bucket.

through the new API makes sure that the core has a reasonable idea of what is going on in the device. At the same time, the notification-based approach is necessary to close the race between the drive report and new hardware datapath flows popping up and making a bucket busy.

Bucket Flags

As individual buckets are offloaded (or configured to trap traffic to software datapath), a driver ought to mark them with either `offload` or `trap` flags.

Testing

A number of selftests have been added in order to check that the algorithm behaves in an expected manner. Specifically for testing edge cases of the bucket migration algorithm, faux `netdevsim` offload has been added.

The module exposes a debugfs interface that allows marking individual buckets as busy. For example, to mark bucket 23 in next-hop group 10 as active, one would write the string "10 23" to file `/sys/kernel/debug/netdevsim/netdevsimX/fib/nexthop_bucket_activity`. Another interface, `.../fib/fail_nexthop_bucket_replace`, allows configuring that the next attempt to migrate a bucket should fail.

These interfaces permit careful testing of the algorithm edge cases. It is for example possible to keep busy the buckets that belong to certain next hops, then lower that next hop's weight, and observe how the algorithm handles the transition.

References

- [1] Hopps, C. 2000. Analysis of an Equal-Cost Multi-Path Algorithm. RFC 2992.

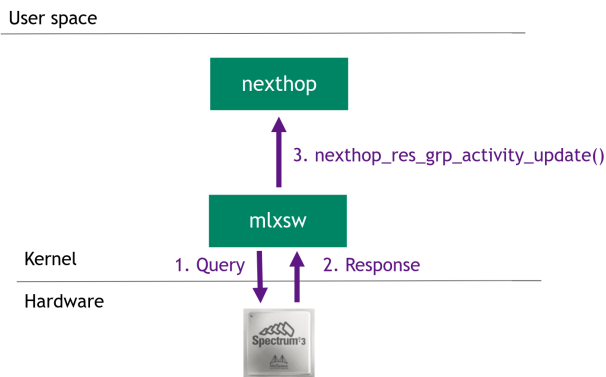


Figure 6: Driver periodically calls `nexthop_res_grp_activity_update()` to keep bucket activity up to date.

This notification-based approach would be enough to implement the raw functionality. However, it would inevitably lead to a lot of ping-pong between the core next-hop code, the driver and the device, as the algorithm would keep proposing migrations that involved busy buckets. Reporting